

Lesson 6: Table Joins

Duration: 20 minutes

Deliverable: `lesson6_joins.sql`

Learning Objectives

By the end of this lesson, you will be able to:

- Understand why we need JOINs
- Write INNER JOIN queries
- Write LEFT JOIN queries
- Combine data from multiple tables
- Handle NULL values in joins
- Use joins with WHERE, ORDER BY, and aggregate functions



What are JOINS?

JOINS combine rows from two or more tables based on a related column. They allow you to retrieve data spread across multiple tables in a single query.

Without JOINs:

- Query characters: Get `homeworld_id = 1`
- Query planets: Find that ID 1 = "Tatooine"
- Manually connect the information

With JOINs:

- One query returns character name AND planet name together!

Types of JOINS

Join Type	Returns	Use When
INNER JOIN	Only matching rows from both tables	You want records that exist in both tables
LEFT JOIN	All rows from left table, matching from right	You want all from table A, even if no match in table B
RIGHT JOIN	All rows from right table, matching from left	SQLite doesn't support; use LEFT JOIN instead

Part 1: INNER JOIN (8 minutes)

Step 1: Create Your SQL File

1. Create: lesson6_joins.sql
2. Add header:

```
-- Lesson 6: Table Joins  
-- Student Name: [Your Name]  
-- Date: [Today's Date]  
--  
-- This script demonstrates INNER and LEFT joins
```

INNER JOIN Syntax

```
SELECT columns  
FROM table1  
INNER JOIN table2 ON table1.column = table2.column;
```

Step 2: Join Characters and Planets

```
-- Query 1: Show characters with their homeworld details  
SELECT  
    characters.name AS character_name,  
    characters.species,  
    planets.name AS homeworld_name,  
    planets.climate  
FROM characters  
INNER JOIN planets ON characters.homeworld_id = planets.id;
```

Explanation:

- characters.name - Specify which table the column comes from
- AS character_name - Rename column in results (avoids confusion when both tables have "name")
- ON characters.homeworld_id = planets.id - How tables relate

Execute this query to see characters with planet details!

Step 3: Using Table Aliases

Table names can be long. Use **aliases** to shorten them:

```
-- Query 2: Same query with table aliases
SELECT
    c.name AS character_name,
    c.species,
    p.name AS planet_name,
    p.climate,
    p.population
FROM characters c
INNER JOIN planets p ON c.homeworld_id = p.id;
```

Much cleaner! `c` is shorthand for `characters`, `p` for `planets`.

Step 4: Join Three Tables

Let's combine characters, vehicles, and the junction table:

```
-- Query 3: Show which characters pilot which vehicles
SELECT
    c.name AS character_name,
    v.name AS vehicle_name,
    v.vehicle_class
FROM characters c
INNER JOIN character_vehicles cv ON c.id = cv.character_id
INNER JOIN vehicles v ON cv.vehicle_id = v.id
ORDER BY c.name;
```

Explanation:

1. Start with `characters` table
2. Join to `character_vehicles` (junction table) linking characters to vehicles
3. Join to `vehicles` table to get vehicle details
4. Order by character name

Step 5: JOIN with WHERE

```
-- Query 4: Find all humans and their homeworlds
SELECT
    c.name,
    c.species,
    p.name AS homeworld
FROM characters c
INNER JOIN planets p ON c.homeworld_id = p.id
WHERE c.species = 'Human';
```

Explanation: JOIN first, then filter with WHERE.

Step 6: JOIN with Aggregate Functions

```
-- Query 5: Count how many characters are from each planet
SELECT
    p.name AS planet_name,
    COUNT(c.id) AS character_count
FROM planets p
INNER JOIN characters c ON p.id = c.homeworld_id
GROUP BY p.name
ORDER BY character_count DESC;
```

Notice: Only shows planets that have characters (INNER JOIN requirement).



Part 2: LEFT JOIN (8 minutes)

LEFT JOIN returns ALL rows from the left table, plus matching rows from the right table. If there's no match, NULL appears for right table columns.

LEFT JOIN Syntax

```
SELECT columns
FROM table1
LEFT JOIN table2 ON table1.column = table2.column;
```

Step 7: Find Characters Without Vehicles

```
-- Query 6: List all characters and their vehicles (including those with no vehicles)
SELECT
    c.name AS character_name,
    v.name AS vehicle_name
FROM characters c
LEFT JOIN character_vehicles cv ON c.id = cv.character_id
LEFT JOIN vehicles v ON cv.vehicle_id = v.id
ORDER BY c.name;
```

Result: Characters without vehicles show NULL for vehicle_name.

Step 8: Filter for Characters WITHOUT Vehicles

```
-- Query 7: Find characters who don't pilot any vehicles
SELECT
    c.name AS character_name,
    c.species
FROM characters c
LEFT JOIN character_vehicles cv ON c.id = cv.character_id
WHERE cv.vehicle_id IS NULL;
```

Explanation: After LEFT JOIN, rows without matches have NULL. We filter for those.

Step 9: Find Vehicles Without Pilots

```
-- Query 8: Find vehicles that no character pilots
SELECT
    v.name AS vehicle_name,
    v.vehicle_class
FROM vehicles v
LEFT JOIN character_vehicles cv ON v.id = cv.vehicle_id
WHERE cv.character_id IS NULL;
```

Step 10: Count Including Empty Groups

```
-- Query 9: Count characters per planet (including planets with 0 characters)
SELECT
    p.name AS planet_name,
    COUNT(c.id) AS character_count
FROM planets p
LEFT JOIN characters c ON p.id = c.homeworld_id
GROUP BY p.name
ORDER BY character_count DESC;
```

Difference from Query 5: This shows ALL planets, even those with no characters (count = 0).

Part 3: Complex JOIN Queries (4 minutes)

Step 11: Multiple JOINS with Filtering

```
-- Query 10: Find humans who pilot starfighters
SELECT
    c.name AS character_name,
    v.name AS vehicle_name,
    v.vehicle_class
FROM characters c
INNER JOIN character_vehicles cv ON c.id = cv.character_id
INNER JOIN vehicles v ON cv.vehicle_id = v.id
WHERE c.species = 'Human' AND v.vehicle_class = 'Starfighter';
```

Step 12: JOIN with Aggregate and HAVING

```
-- Query 11: Find characters who pilot more than one vehicle
SELECT
    c.name AS character_name,
    COUNT(v.id) AS vehicle_count
FROM characters c
INNER JOIN character_vehicles cv ON c.id = cv.character_id
INNER JOIN vehicles v ON cv.vehicle_id = v.id
GROUP BY c.name
HAVING COUNT(v.id) > 1;
```

Step 13: Comprehensive Query

```
-- Query 12: Character summary with all related data
SELECT
    c.name AS character,
    c.species,
    p.name AS homeworld,
    p.climate,
    COUNT(v.id) AS vehicles_piloted
FROM characters c
LEFT JOIN planets p ON c.homeworld_id = p.id
LEFT JOIN character_vehicles cv ON c.id = cv.character_id
LEFT JOIN vehicles v ON cv.vehicle_id = v.id
GROUP BY c.name, c.species, p.name, p.climate
ORDER BY c.name;
```

Practice Exercises

Exercise 1: Simple INNER JOIN

```
-- Exercise 1: List all characters with their homeworld's population
SELECT
    c.name,
    p.name AS homeworld,
    p.population
FROM characters c
INNER JOIN planets p ON c.homeworld_id = p.id;
```

Exercise 2: Multiple JOINS

```
-- Exercise 2: Show all vehicle-pilot pairs with character species
SELECT
    c.name AS pilot,
    c.species,
    v.name AS vehicle
FROM characters c
INNER JOIN character_vehicles cv ON c.id = cv.character_id
INNER JOIN vehicles v ON cv.vehicle_id = v.id;
```

Exercise 3: LEFT JOIN with NULL Check

```
-- Exercise 3: Find all planets with no characters
SELECT
    p.name AS planet_name,
    p.climate
FROM planets p
LEFT JOIN characters c ON p.id = c.homeworld_id
WHERE c.id IS NULL;
```

Exercise 4: Aggregate with JOIN

```
-- Exercise 4: Show each vehicle with the count of who pilots it
SELECT
    v.name AS vehicle_name,
    COUNT(c.id) AS pilot_count
FROM vehicles v
LEFT JOIN character_vehicles cv ON v.id = cv.vehicle_id
LEFT JOIN characters c ON cv.character_id = c.id
GROUP BY v.name;
```



Common Errors & Troubleshooting

Error: "ambiguous column name"

Problem: Column exists in multiple tables and you didn't specify which.

Wrong:

```
SELECT name FROM characters
INNER JOIN planets ON homeworld_id = id;
```

Correct:

```
SELECT characters.name FROM characters
INNER JOIN planets ON characters.homeworld_id = planets.id;
```

Error: "no such column"

Problem: Misspelt column or using wrong table prefix.

Solution: Verify column names:

```
-- Check column names in each table
PRAGMA table_info(characters);
PRAGMA table_info(planets);
```

Wrong JOIN Type

Symptom: Missing expected rows.

Problem: Used INNER JOIN when you needed LEFT JOIN.

Remember:

- INNER JOIN: Only rows with matches in BOTH tables
- LEFT JOIN: ALL rows from left table, matches from right

Incorrect ON Clause

Problem: Joining on wrong columns.

Wrong:

```
FROM characters c
INNER JOIN planets p ON c.id = p.id -- Wrong columns!
```

Correct:

```
FROM characters c
INNER JOIN planets p ON c.homeworld_id = p.id -- Correct relationship
```

Cartesian Product (Too Many Results)

Problem: Missing ON clause creates every possible combination.

Wrong:

```
SELECT * FROM characters, planets; -- Returns 11 × 8 = 88 rows!
```

Correct:

```
SELECT * FROM characters
INNER JOIN planets ON characters.homeworld_id = planets.id;
```

Checkpoint: What You've Learnt

Before moving on, make sure you can:

-  Explain the purpose of JOINs
-  Write INNER JOIN queries
-  Write LEFT JOIN queries
-  Use table aliases for clarity
-  Join three or more tables
-  Combine JOINs with WHERE, GROUP BY, and HAVING
-  Find records with no matches using LEFT JOIN and NULL

Challenge Problem (Optional)

Task: Create a query that shows each planet with:

- Planet name
- Climate
- Number of characters from that planet
- Average height of characters from that planet
- Include planets with NO characters (show 0 for count, NULL for average)
- Order by character count (descending), then planet name

Hints:

- Use LEFT JOIN
- Use COUNT() and AVG()
- Use GROUP BY
- Use ORDER BY with multiple columns

Click to reveal the solution

```
SELECT
    p.name AS planet_name,
    p.climate,
    COUNT(c.id) AS character_count,
    ROUND(AVG(c.height), 1) AS avg_height
FROM planets p
LEFT JOIN characters c ON p.id = c.homeworld_id
GROUP BY p.name, p.climate
ORDER BY character_count DESC, p.name;
```

Save Your Work with Git

```
git status  
git add lessons/lesson6_joins.sql database/starwars.db  
git commit -m "Completed Lesson 6: INNER and LEFT joins across multiple tables"  
git push
```



Key SQL Commands Learnt

Command	Purpose	Example
INNER JOIN	Return only matching rows	INNER JOIN planets ON c.homeworld_id = p.id
LEFT JOIN	Return all from left table	LEFT JOIN vehicles ON c.id = v.pilot_id
ON	Specify join condition	ON table1.id = table2.foreign_id
Table Alias	Shorten table names	FROM characters c
IS NULL	Check for NULL values	WHERE cv.vehicle_id IS NULL



JOIN Visual Reference

INNER JOIN:

Table A: [1, 2, 3]

Table B: [2, 3, 4]

Result: [2, 3] ← Only matches

LEFT JOIN:

Table A: [1, 2, 3]

Table B: [2, 3, 4]

Result: [1, 2, 3] ← All from A, with B data where available



Fantastic Work!

You can now combine data from multiple tables! In the next lesson, you'll learn how to modify data with UPDATE and DELETE statements.

Ready to continue? Move on to `lesson7_instructions.md`

Need Help?

- Draw table relationships on paper
- Test JOIN without WHERE first
- Check which table has which columns
- Verify foreign key relationships
- Ask your instructor!